

# Implementacja sieci neuronowej

```
In [22]: def getModel():
model = tf.keras.models.Sequential([
    tf.keras.layers.Conv2D(8,(3,3), activation='relu',input_shape=(28, 28,1)),
    tf.keras.layers.Conv2D(16,(3,3), activation='relu'),
    tf.keras.layers.Conv2D(32,(3,3), activation='relu'),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(10, activation='softmax')])
model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])
return model
model = getModel()
model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 26, 26, 8)	80
conv2d_1 (Conv2D)	(None, 24, 24, 16)	1168
conv2d_2 (Conv2D)	(None, 22, 22, 32)	4640
flatten (Flatten)	(None, 15488)	0
dense (Dense)	(None, 10)	154890


=====  
Total params: 160,778  
Trainable params: 160,778  
Non-trainable params: 0  
=====

```
In [24]: epochs=10
batch_size=256

history = model.fit(X_train, y_train, epochs=epochs, batch_size=batch_size,
                    validation_data=(X_val,y_val),)
```


```
In [25]: plt.figure(figsize=(20,10))
plt.plot(history.history['accuracy'], label='acc')
plt.plot(history.history['val_accuracy'], label='val_acc')
plt.legend()
```

Out[25]: <matplotlib.legend.Legend at 0x12beac56310>




# Optymalizacja – zmiana funkcji aktywacji


## Relu + Adam




## Swish + Adam



## Relu + RMSprop



## Swish + RMSprop






# Optymalizacja - Dropout

```
In [26]: def getModelDropout(ratios=[0.0,0.0]):  
    model = tf.keras.models.Sequential([  
        tf.keras.layers.Conv2D(8,(3,3), activation='relu',input_shape=(28, 28,1)),  
        tf.keras.layers.Dropout(ratios[0]),  
        tf.keras.layers.Conv2D(32,(3,3), activation='relu'),  
        tf.keras.layers.Flatten(),  
        tf.keras.layers.Dropout(ratios[1]),  
        tf.keras.layers.Dense(10, activation='softmax')])  
    model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=0.001),  
        loss='sparse_categorical_crossentropy',  
        metrics=['accuracy'])  
    return model  
model = getModelDropout([0,0])  
model.summary()  
#del model
```

Model: "sequential\_8"

Layer (type)	Output Shape	Param #
conv2d_18 (Conv2D)	(None, 26, 26, 8)	80
dropout_4 (Dropout)	(None, 26, 26, 8)	0
conv2d_19 (Conv2D)	(None, 24, 24, 32)	2336
flatten_8 (Flatten)	(None, 18432)	0
dropout_5 (Dropout)	(None, 18432)	0
dense_8 (Dense)	(None, 10)	184330
Total params: 186,746		
Trainable params: 186,746		
Non-trainable params: 0		

Dropout



# Optymalizacja – Keras Tuner

```
In [6]: from keras_tuner import RandomSearch

def getModel(hp):
    l1_l2=tf.keras.regularizers.L1_L2(hp.Float('l1_l2', 1e-8, 0.01, sampling='log'))
    l=[tf.keras.layers.Conv2D(8,(3,3), padding='same', activations='relu',
                               kernel_regularizer=l1_l2, input_shape=(img_size, img_size, 3)),]
    pooling = hp.Int('avg_pooling',0,3)
    bn=hp.Int('BN',0,2)

    if bn>0:
        l.append(tf.keras.layers.BatchNormalization())
    for i in range(hp.Int('convs',0,4)):
        if pooling==1:
            l.append(tf.keras.layers.AvgPool2D())
        if pooling==2:
            l.append(tf.keras.layers.MaxPool2D())
        l.append(tf.keras.layers.Conv2D(2**(5+i),(3,3), padding='same', activation='relu',
                                         kernel_regularizer=l1_l2))
    if bn==2:
        l.append(tf.keras.layers.BatchNormalization())
        l.append(tf.keras.layers.GlobalAveragePooling2D())
        l.append(tf.keras.layers.Dropout(hp.Float('dropout',0,0.9)))
        l.append(tf.keras.layers.Dense(num_classes, activation='softmax'))

    model = tf.keras.models.Sequential(l)
    model.compile(optimizer=tf.keras.optimizers.Adam(hp.Float('lr',0.000001, 0.01, sampling='log')),
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])

    return model

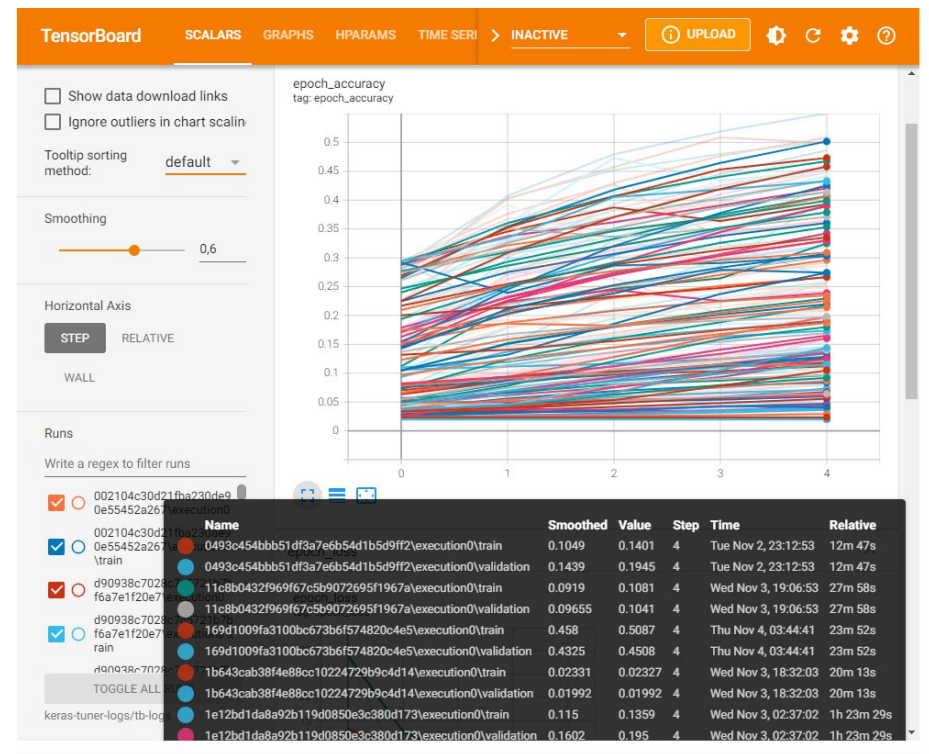
In [7]: tuner =RandomSearch(getModel,
                             objective="val_accuracy",
                             max_trials=15,
                             executions_per_trial=3, # Uśrednianie wyników
                             overwrite=True,
                             directory="ktuner",
                             project_name="resisc45")

In [9]: tuner.search_space_summary()
```

Search space summary  
Default search space size: 6  
l2 (Float)  
{'default': 1e-08, 'conditions': [], 'min\_value': 1e-08, 'max\_value': 0.01, 'step': None, 'sampling': 'log'}  
avg\_pooling (Int)  
{'default': None, 'conditions': [], 'min\_value': 0, 'max\_value': 2, 'step': 1, 'sampling': None}  
BN (Int)  
{'default': None, 'conditions': [], 'min\_value': 0, 'max\_value': 2, 'step': 1, 'sampling': None}  
convs (Int)  
{'default': None, 'conditions': [], 'min\_value': 0, 'max\_value': 3, 'step': 1, 'sampling': None}  
dropout (Float)  
{'default': 0.0, 'conditions': [], 'min\_value': 0.0, 'max\_value': 0.9, 'step': None, 'sampling': None}  
lr (Float)  
{'default': 1e-05, 'conditions': [], 'min\_value': 1e-05, 'max\_value': 0.1, 'step': None, 'sampling': 'log'}

```
In [1]: from torch.utils import tensorboard
```

```
In [2]: %reload_ext tensorboard
%tensorboard --logdir keras-tuner-logs/tb-logs --port 6001
```



# Optymalizacja – Keras Tuner



```
In [13]: models[0].summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 256, 256, 8)	224
batch_normalization (Batch Normalization)	(None, 256, 256, 8)	32
max_pooling2d (MaxPooling2D)	(None, 128, 128, 8)	0
conv2d_1 (Conv2D)	(None, 128, 128, 32)	2336
batch_normalization_1 (Batch Normalization)	(None, 128, 128, 32)	128
global_average_pooling2d (Global Average Pooling2D)	(None, 32)	0
dropout (Dropout)	(None, 32)	0
dense (Dense)	(None, 45)	1485
		-----
Total params:		4,205
Trainable params:		4,125
Non-trainable params:		80

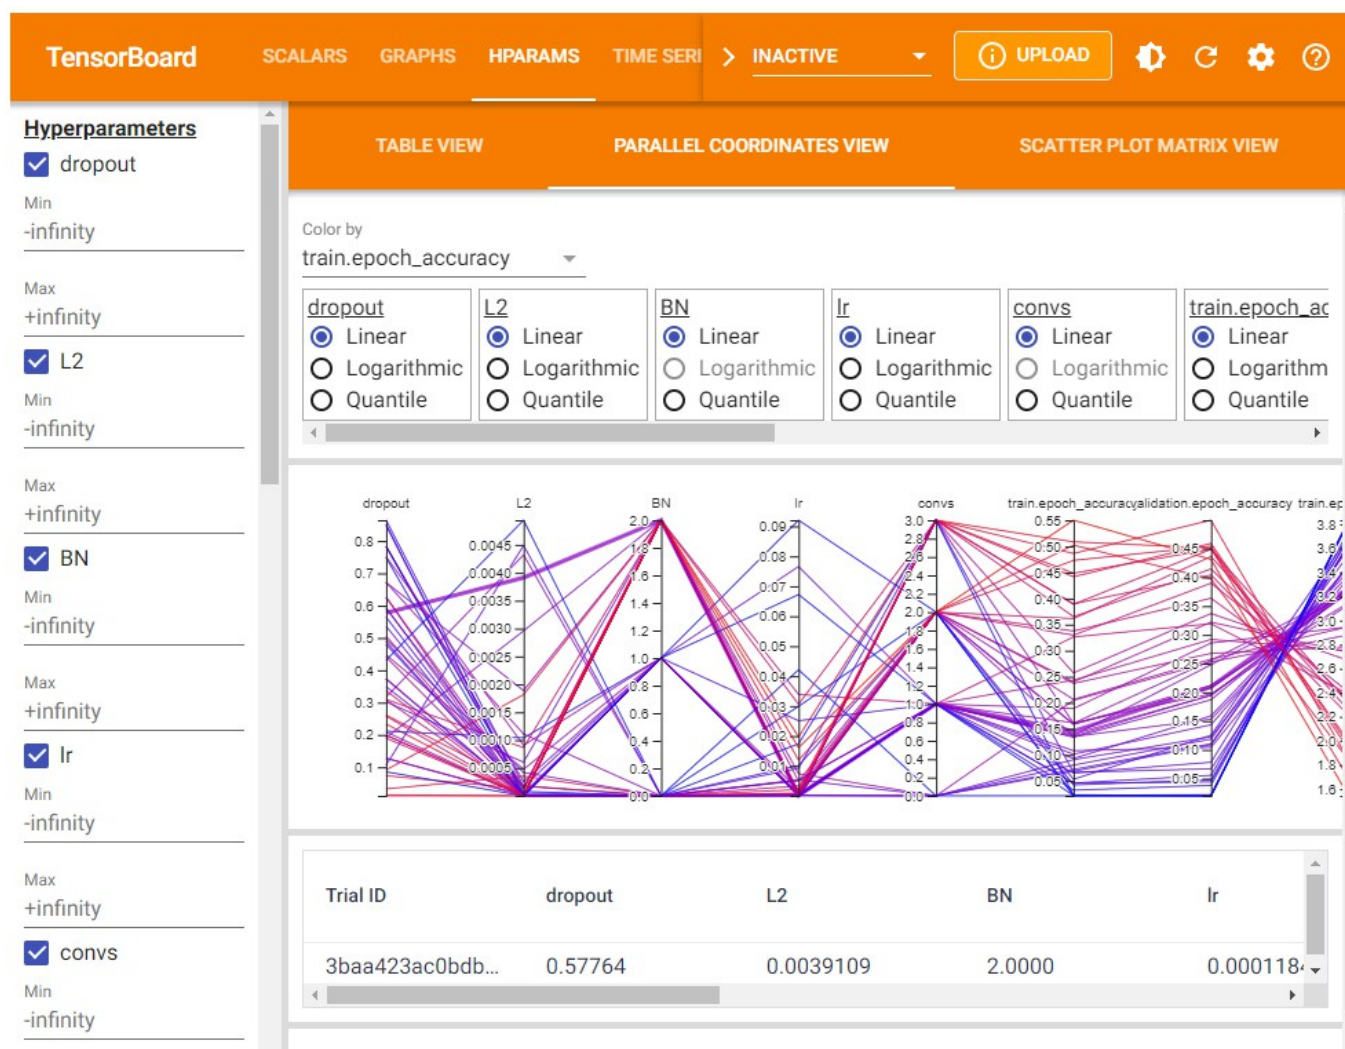
```
In [14]: models[1].summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 256, 256, 8)	224
batch_normalization (Batch Normalization)	(None, 256, 256, 8)	32
conv2d_1 (Conv2D)	(None, 256, 256, 32)	2336
conv2d_2 (Conv2D)	(None, 256, 256, 64)	18496
conv2d_3 (Conv2D)	(None, 256, 256, 128)	73856
global_average_pooling2d (Global Average Pooling2D)	(None, 128)	0
dropout (Dropout)	(None, 128)	0
dense (Dense)	(None, 45)	5805
		-----
Total params:		100,749
Trainable params:		100,733
Non-trainable params:		16



# Optymalizacja – Keras Tuner



# Optymalizacja - Optuna

```
In [6]: import optuna
from optuna.trial import TrialState
from optuna.integration import TFKerasPruningCallback, TensorBoardCallback

def getModel(lr, pooling, flatten, convs):
    l = [tf.keras.layers.Conv2D(8,(3,3), padding='same', activation='relu',input_shape=(img_size, img_size,3)),]
    for i in range(int(convs)):
        if pooling=='avg':
            l.append(tf.keras.layers.AvgPool2D())
        if pooling=='max':
            l.append(tf.keras.layers.MaxPool2D())
        l.append(tf.keras.layers.Conv2D(2**(i+4),(3,3), padding='same', activation='relu'))
    if flatten=='global':
        l.append(tf.keras.layers.GlobalAveragePooling2D())
    else:
        l.append(tf.keras.layers.Flatten())
    l.append(tf.keras.layers.Dense(num_classes, activation='softmax'))

    model = tf.keras.models.Sequential(l)
    model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=lr),
                  loss='sparse_categorical_crossentropy',
                  metrics=['accuracy'])

    return model

params={'lr':0.001,
        'pooling':'avg',
        'flatten':'global',
        'convs':5}

model = getModel(**params)
model.summary()
del model
```

Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 256, 256, 8)	224
average_pooling2d (AveragePo	(None, 128, 128, 8)	0
conv2d_1 (Conv2D)	(None, 128, 128, 16)	1168
average_pooling2d_1 (Average	(None, 64, 64, 16)	0
conv2d_2 (Conv2D)	(None, 64, 64, 32)	4640
average_pooling2d_2 (Average	(None, 32, 32, 32)	0
conv2d_3 (Conv2D)	(None, 32, 32, 64)	18496
average_pooling2d_3 (Average	(None, 16, 16, 64)	0
conv2d_4 (Conv2D)	(None, 16, 16, 128)	73856
average_pooling2d_4 (Average	(None, 8, 8, 128)	0
conv2d_5 (Conv2D)	(None, 8, 8, 256)	295168
global_average_pooling2d (G1	(None, 256)	0
dense (Dense)	(None, 45)	11565
-----		
Total params: 405,117		
Trainable params: 405,117		
Non-trainable params: 0		

```
In [7]: epochs = 20
```

```
In [8]: def objective(trial):
    convs = trial.suggest_int("convs", 0, 3)
    pooling = trial.suggest_categorical("pooling", ["no", "avg", "max"])
    flatten = trial.suggest_categorical("flatten", ["global", "flatten"])
    lr = trial.suggest_float("lr", 1e-5, 1e-1, log=True)

    model = getModel(lr, pooling, flatten, convs)

    history = model.fit(ds_train, # 10% of data
                        epochs=epochs,
                        batch_size=batch_size,
                        verbose=0,
                        validation_freq=1,
                        validation_data=(ds_val),
                        callbacks=[TFKerasPruningCallback(trial, "val_accuracy")])

    return history.history['val_accuracy'][-1] # return last val acc
```

```
In [9]: study = optuna.create_study(direction="maximize") # --> maximize accuracy
study.optimize(objective, n_trials=100, timeout=None,
               callbacks=[TensorBoardCallback('optuna-keras-logs2', 'val_accuracy')])

pruned_trials = study.get_trials(deepcopy=False, states=[TrialState.PRUNED])
complete_trials = study.get_trials(deepcopy=False, states=[TrialState.COMPLETE])

print("Study statistics: ")
print("  Number of finished trials: ", len(study.trials))
print("  Number of pruned trials: ", len(pruned_trials))
print("  Number of complete trials: ", len(complete_trials))

print("Best trial:")
trial = study.best_trial

print("  Value: ", trial.value)

print("  Params: ")
for key, value in trial.params.items():
    print("    {}: {}".format(key, value))
```

# Optymalizacja - Optuna

Tooltip sorting method: default

Smoothing 0,6

Horizontal Axis


STEP RELATIVE

WALL

Runs

Write a regex to filter runs

- ☒ trial-63
- ☒ trial-78
- ☒ trial-82
- ☒ trial-85
- ☒ trial-91
- ☒ trial-92
- ☒ trial-93



```
atten': 'global', 'lr': 0.0016557650763173051}. Best is trial 11 with value: 0.6713982820510864.
[I 2021-12-02 08:48:41,387] Trial 62 finished with value: 0.6537076234617505 and parameters: {'conv': 3, 'pooling': 'avg', 'fl
atten': 'global', 'lr': 0.0014974906870220015}. Best is trial 11 with value: 0.6713982820510864.
[I 2021-12-02 09:57:26,664] Trial 63 finished with value: 0.6447033882141113 and parameters: {'conv': 3, 'pooling': 'avg', 'fl
atten': 'global', 'lr': 0.0017040832215024272}. Best is trial 11 with value: 0.6713982820510864.
[I 2021-12-02 10:04:20,063] Trial 64 pruned. Trial was pruned at epoch 1.
[I 2021-12-02 10:07:46,105] Trial 65 pruned. Trial was pruned at epoch 0.
[I 2021-12-02 10:41:51,409] Trial 66 pruned. Trial was pruned at epoch 9.
[I 2021-12-02 11:20:00,827] Trial 67 pruned. Trial was pruned at epoch 10.
[I 2021-12-02 11:22:47,409] Trial 68 pruned. Trial was pruned at epoch 0.
[I 2021-12-02 11:53:45,901] Trial 69 pruned. Trial was pruned at epoch 0.
[I 2021-12-02 11:57:14,581] Trial 70 pruned. Trial was pruned at epoch 0.
[I 2021-12-02 12:00:43,116] Trial 71 pruned. Trial was pruned at epoch 0.
[I 2021-12-02 12:27:54,834] Trial 72 pruned. Trial was pruned at epoch 7.
[I 2021-12-02 12:31:17,913] Trial 73 pruned. Trial was pruned at epoch 0.
[I 2021-12-02 12:34:37,388] Trial 74 pruned. Trial was pruned at epoch 0.
[I 2021-12-02 12:50:52,367] Trial 75 pruned. Trial was pruned at epoch 4.
[I 2021-12-02 12:57:25,822] Trial 76 pruned. Trial was pruned at epoch 1.
[I 2021-12-02 13:00:47,367] Trial 77 pruned. Trial was pruned at epoch 0.
[I 2021-12-02 14:05:32,490] Trial 78 finished with value: 0.6325212121009827 and parameters: {'conv': 3, 'pooling': 'avg', 'fl
atten': 'global', 'lr': 0.001156945889121495}. Best is trial 11 with value: 0.6713982820510864.
[I 2021-12-02 14:08:48,621] Trial 79 pruned. Trial was pruned at epoch 0.
[I 2021-12-02 14:11:47,315] Trial 80 pruned. Trial was pruned at epoch 0.
[I 2021-12-02 14:15:04,002] Trial 81 pruned. Trial was pruned at epoch 0.
[I 2021-12-02 15:20:33,607] Trial 82 finished with value: 0.6525423526763916 and parameters: {'conv': 3, 'pooling': 'avg', 'fl
atten': 'global', 'lr': 0.0015587859065636276}. Best is trial 11 with value: 0.6713982820510864.
[I 2021-12-02 15:23:51,562] Trial 83 pruned. Trial was pruned at epoch 0.
[I 2021-12-02 15:30:19,544] Trial 84 pruned. Trial was pruned at epoch 1.
[I 2021-12-02 16:34:43,643] Trial 85 finished with value: 0.6709745526313782 and parameters: {'conv': 3, 'pooling': 'avg', 'fl
atten': 'global', 'lr': 0.0027921673799563237}. Best is trial 11 with value: 0.6713982820510864.
[I 2021-12-02 16:38:01,528] Trial 86 pruned. Trial was pruned at epoch 0.
[I 2021-12-02 17:00:05,823] Trial 87 pruned. Trial was pruned at epoch 0.
[I 2021-12-02 17:11:03,914] Trial 88 pruned. Trial was pruned at epoch 2.
[I 2021-12-02 17:14:27,161] Trial 89 pruned. Trial was pruned at epoch 0.
[I 2021-12-02 17:17:49,497] Trial 90 pruned. Trial was pruned at epoch 0.
[I 2021-12-02 18:24:58,778] Trial 91 finished with value: 0.6551907062530518 and parameters: {'conv': 3, 'pooling': 'avg', 'fl
atten': 'global', 'lr': 0.0020396384183621755}. Best is trial 11 with value: 0.6713982820510864.
[I 2021-12-02 19:32:01,470] Trial 92 finished with value: 0.6605932116500484 and parameters: {'conv': 3, 'pooling': 'avg', 'fl
atten': 'global', 'lr': 0.0019447596710891887}. Best is trial 11 with value: 0.6713982820510864.
[I 2021-12-02 20:40:31,640] Trial 93 finished with value: 0.6691737174907793 and parameters: {'conv': 3, 'pooling': 'avg', 'fl
atten': 'global', 'lr': 0.0020273400701534407}. Best is trial 11 with value: 0.6713982820510864.
[I 2021-12-02 21:01:13,136] Trial 94 pruned. Trial was pruned at epoch 5.
[I 2021-12-02 21:14:33,838] Trial 95 pruned. Trial was pruned at epoch 3.
[I 2021-12-02 21:44:30,326] Trial 96 pruned. Trial was pruned at epoch 8.
[I 2021-12-02 21:47:43,820] Trial 97 pruned. Trial was pruned at epoch 0.
[I 2021-12-02 21:50:16,466] Trial 98 pruned. Trial was pruned at epoch 0.
[I 2021-12-02 21:53:35,047] Trial 99 pruned. Trial was pruned at epoch 0.
```

Study statistics:  
Number of finished trials: 100  
Number of pruned trials: 67  
Number of complete trials: 33  
Best trial:  
Value: 0.6713982820510864  
Params:  
conv: 3  
pooling: avg  
flatten: global  
lr: 0.0013234805544161227




# Augumentacja – gdy mamy za mało danych

bez augmentacji

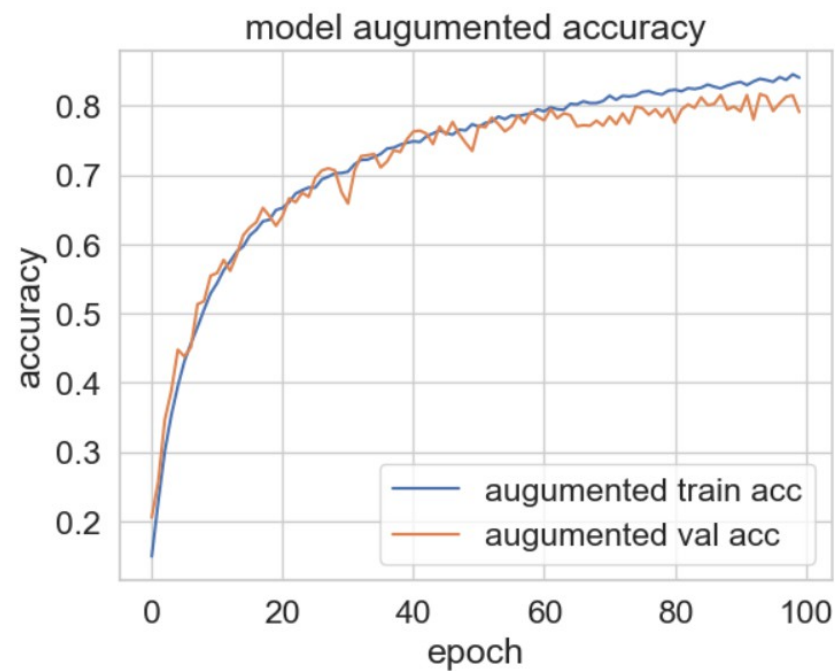



z augmentacją



```
In [25]: data_augmentation = tf.keras.Sequential([tf.keras.layers.RandomContrast(factor=0.2),  
tf.keras.layers.RandomRotation(factor=0.1),  
tf.keras.layers.RandomFlip('horizontal'),  
tf.keras.layers.RandomFlip('vertical'),  
tf.keras.layers.RandomTranslation(height_factor=0.2, width_factor=0.2),  
tf.keras.layers.RandomZoom(height_factor=0.5, width_factor=0.2)])
```

# Augumentacja - wyniki

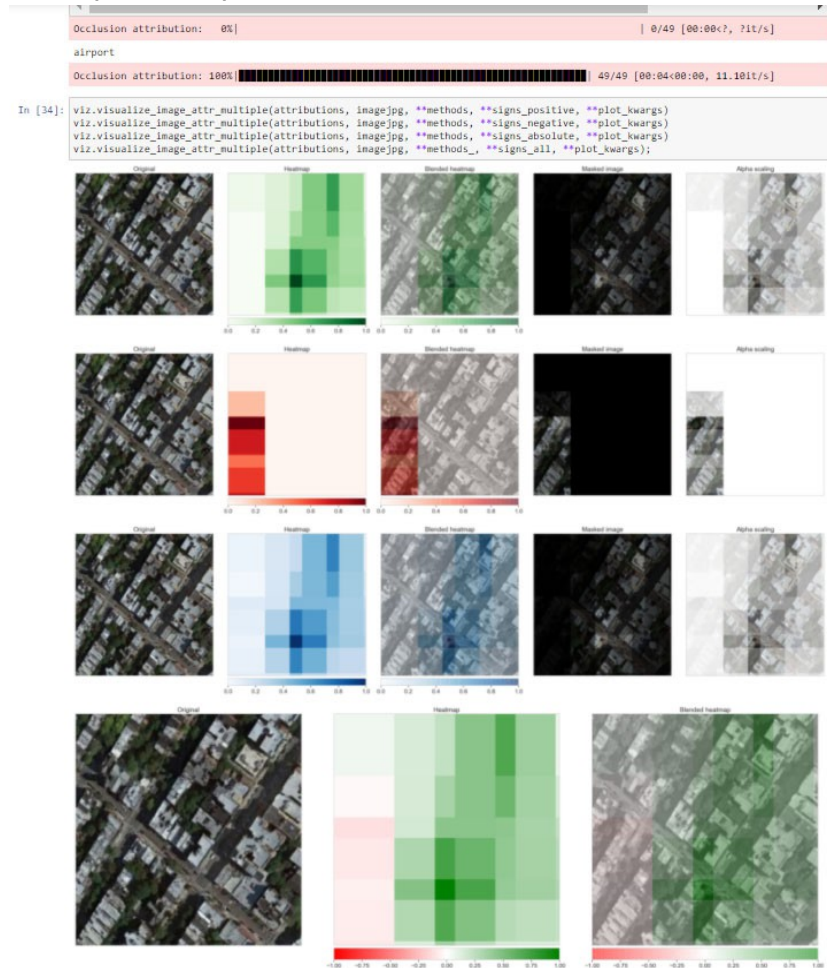


# Klasyfikacja obrazu – analiza modelu (pretrained model resnet50)

- Occlusion sensitivity – klasa właściwa (dense residential)



- Occlusion sensitivity – klasa niewłaściwa (airport)





# Klasyfikacja obrazu – analiza modelu (pretrained model resnet50)

- (Guided) GradCAM – klasa właściwa (dense residential)
- (Guided) GradCAM – klasa niewłaściwa (golf course)






# Klasyfikacja obrazu - Niezbyt udane wdrożenie jako aplikacja demo czyli jak oszukać AI (gradio) ;)

Predykcja dla zbioru stl10

wizualizacja

IM



OUTPUT

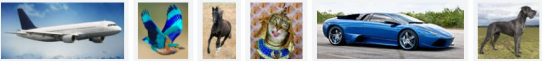
horse

horse	100%
truck	0%
monkey	0%
bird	0%
cat	0%
dog	0%
airplane	0%
car	0%
deer	0%
ship	0%

Clear Submit

Screenshot Flag

Examples

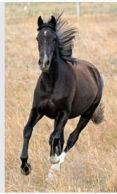


[view the api](#) • [built with](#)

Predykcja dla zbioru stl10

wizualizacja

IM



OUTPUT

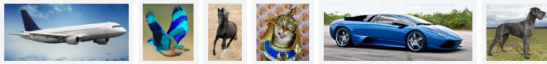
horse

horse	100%
dog	0%
monkey	0%
car	0%
truck	0%
airplane	0%
bird	0%
cat	0%
deer	0%
ship	0%

Clear Submit

Screenshot Flag

Examples




[view the api](#) • [built with](#)

# Gradio kontra pretrenowane modele ;)

## MobileNetV2 vs EfficientNetB0

Porównanie dwóch modeli z TF Keras

IM



MOBILENET PREDICT: OUTPUT 0.05s

**plate**

plate	29%
wok	10%
frying pan	7%
hot pot	5%
chocolate s...	3%

EFFNET PREDICT: OUTPUT

**carbonara**

carbonara	21%
plate	10%
hot pot	7%
pizza	4%
frying pan	4%

Clear Submit

Screenshot Flag

Examples



